



Vrije Universiteit Brussel

Faculteit Wetenschappen
Departement Informatica
en Toegepaste Informatica

Content Migration and Layout for the MindXpres Presentation Tool

Proefschrift ingediend met het oog op het behalen
van de graad van Master in de Toegepaste Informatica

Joris Vandermeersch

Promotor: Prof. Dr. Beat Signer
Begeleider: Reinout Roels

Augustus 2015





Vrije Universiteit Brussel

Faculty of Science
Department of Computer Science
and Applied Computer Science

Content Migration and Layout for the MindXpres Presentation Tool

Graduation thesis submitted in partial fulfillment of the
requirements for the degree of Master in Applied Computer Science

Joris Vandermeersch

Promotor: Prof. Dr. Beat Signer
Advisor: Reinout Roels

August 2015



ABSTRACT

Microsoft PowerPoint continues to be used worldwide in staggering numbers. We try to provide an alternative with MindXpres, facilitating the switch by converting existing PowerPoint presentations into MindXpres presentations, and automatically fixing the layout in the process.

ACKNOWLEDGEMENTS

*“Simplicity is a great virtue,
but it requires hard work to achieve it and education to appreciate it.
And to make matters worse: complexity sells better.”*
— Edsger W. Dijkstra

CONTENTS

1. <i>Introduction</i>	6
2. <i>Problem statement</i>	9
2.1 Terminology	10
2.2 Stuff we want to solve	11
2.2.1 Real-life slideware problems	11
3. <i>Related work</i>	13
3.1 Background	13
3.2 Existing solutions	15
3.3 New solutions	18
3.4 MindXpres Platform	20
3.4.1 Document Format and Authoring Language	20
3.4.2 Compiler	21
3.4.3 MindXpres Presentation Bundle	21
3.4.4 Plug-in Types	22
3.4.5 Implementation	23
3.4.6 Use Cases	26
3.4.7 Discussion and Future Work	29
3.4.8 Conclusion	30
3.5 Layout	31
3.5.1 Algorithms	33
4. <i>Approach</i>	34
4.1 Conversion process	35
4.2 Compiler optimizations	37
4.3 Using MindXpres	38
4.3.1 A MindXpres plug-in	39
4.3.2 An automated layout algorithm	40
5. <i>Implementation</i>	42
5.1 Taking PowerPoint apart	43
5.1.1 Bullets	44

5.1.2	Animations	44
5.2	Generating MindXpres	45
5.2.1	Playing MindXpres compiler	45
5.2.2	MindXpres XML	46
5.3	Creating layouts	46
5.3.1	Using constraints	46
5.3.2	Other ways	46
6.	<i>Conclusions and Future Work</i>	47
6.1	Contribution	47
6.2	Future Work	47

1. INTRODUCTION

For over 25 years, Microsoft PowerPoint has been the market leader in digital presentations. Admittedly, it was a revolutionary software package when it was first introduced, and its ease-of-use combined with its supreme graphical capabilities — at least compared to other software in the same era — quickly made it one of the most popular software packages in history. 25 years later, Microsoft PowerPoint can claim over 90% market share in presentation software, and on average 30 million PowerPoint presentations are created every day.

In this time, Microsoft PowerPoint has gotten many new features, and certainly improved and grew with every new version, but it never really changed its core approach. It started out mimicking the then-popular and widespread use of dia and overhead projection slides, which was at the time a good way to convince people of its purpose, allowing them to feel comfortable with a familiar format instead of alienating potential customers with a new and potentially confusing interface.

However, this interface is quite restricting, and in recent years different approaches have seen the light of day. The zoomable user interface of Prezi is probably the most well-known, but apart from abandoning the traditional slide format it does little to improve or extend the concept of presenting information to an audience.

This is where MindXpres comes in. Its extensible plug-in system allows anyone with some knowledge of programming to create new functionality to use in presentations. Examples are interactivity with the audiencer through various means, controlling the presentation from another device — or several! — and (re)modelling data while presenting it, based on feedback from the audience.

While this is obviously a big improvement on the traditional presentation model of Microsoft PowerPoint and the likes, it remains hard to convince the general public of its merits. People are generally afraid of change, and it is important to make the transition as smooth as possible. On top of that, people are often worried that the work they did in the past may be lost — or worse, irrelevant — after switching to something new. This alone may be a huge factor in deciding whether or not to start using new software, or to stick

with what they know.

That is where the subject of this thesis comes in. We aim to provide a way for people to convert their existing PowerPoint presentations into MindXpres presentations, allowing them to take their previous work with them in their switch to MindXpres. This way, we lower the threshold for them to make the decision to start using MindXpres as their presentation software of choice. Once all their existing PowerPoint content is available, usable and editable in MindXpres, it should be obvious to anyone why MindXpres is the better option for their presentations.

Another common problem with PowerPoint presentations is the way they look. This is not necessarily the fault of the software; most people just are not trained in graphical design, and as such they know very little about proper layout, color choices, or slide content limits. Everyone has probably encountered slides with full paragraphs of text, too small to read and / or too much to process in the short time the slide is visible — (too) many people have made those slides themselves.

When we say this is not the fault of the software, that is mostly true, as the creators of these slides obviously made a conscious choice to make their content appear like that. It could be said however that Microsoft PowerPoint and other presentation tools are guilty through inaction. We believe it is possible to have software either warn its users against these choices and practices, or — even better — have the software fix these problems automatically.

One of the primary purposes of MindXpres is to provide automated layout, much like \LaTeX does, ensuring that the content creator only has to worry about the actual content, while the software takes care of layout. In practice, both \LaTeX and MindXpres currently use template-based layouts, where the contents' position is predefined in the template and not related to or based on its size, shape or nature. In the end, everyone who has ever used \LaTeX knows that sooner or later you will struggle to get a certain image incorporated in the text correctly, ending up doing the layout yourself anyway, because the predefined template just doesn't work properly for your specific content.

Our goal is to eradicate those situations. Automated layout should dynamically adjust to any content it is given, no matter the size or aspect ratio. This may seem hard, if you consider the limits of slides and the fact that you can only fit so much content on them before they are full. This is where another important aspect of MindXpres comes into play: we are not necessarily bound to the limits of slides. If we don't have to consider the boundaries of traditional slides, we can fit content together in an aesthetically pleasing way much easier, without having to scale anything.

As such, the second part of this thesis focuses on implementing true

automated layout in MindXpres. Again primarily to convince Microsoft PowerPoint users to switch, showing that their presentations actually could look better in MindXpres, while thus also providing new functionality to existing MindXpres users.

2. PROBLEM STATEMENT

Computers, software and digital content are everywhere. Everything we use nowadays is somehow related to computers and electronics, and if it isn't, it probably will be soon. This may be a bit of hyperbole, but there's a core of truth in it. If you think about it, more and more things have become and are becoming some kind of computer. Coffee machines used to be simple machines that heated water and let it drip over coffee grounds; now there are coffee machines that are connected to the internet, and can be turned on remotely from your smartphone. That smartphone itself is an incredible evolution as well: just 20 years ago, phones were analog devices, and you could use them to call people and nothing more. Today, our phone does a lot more than that, so much more that calling has actually become a minor feature to most people.

Content is going the same way. Photos used to be on a special film, and could be 'developed' onto special paper through a process involving a dark room and several chemicals. Movies existed on a projection film, newspapers were actually made of paper and music was available on vinyl disks with grooves that matched the sound waves. All of this content has been digitized since. This means of course that you can see or hear it using a computer, like you would've seen it without a computer before, but on top of that it means the content can be much more dynamic. You can link it to other content, you can make it respond to your actions, you can discuss it with people around the world. Digitized content allows for interactivity, so that the audience is no longer a passive onlooker but an active participant.

It is no surprise, then, that slideshow presentations have evolved from the original dia's or overhead projection slides into a digital form as well. Except, until recently the evolution stopped there. Slideshows did not become interactive, and the audience remained passive onlookers watching a series of images projected on a screen or a wall. The presenter told a story, and the audience listened. Often during or at the end of the presentation there would be a chance to ask questions, but those questions could only be answered vocally by the presenter. If the question needed any visual explanation, the slideshow would not be able to help. We had digital slides, but the difference with the physical slides was neglectable.

In our eyes, the culprit for this is Microsoft's PowerPoint. This software package took the world by storm, making it possible for everyone with a computer to make digital slideshows, which was impressive at the time. However, Microsoft PowerPoint never really evolved beyond that, except for adding some simple animations allowing text to 'fly' into view (among others). Since it was — and still is! — the dominant player in the world of slideware with over 90% market share, this apathy towards change firmly rooted slideware in the concepts of the past. Luckily, a few years ago some people realized this and decided to take matters into their own hands. They stepped away from the classic slide format, allowing for any kind of layout, combined with zoomable interfaces and other methods of displaying data. One such alternative is MindXpres, created in the WISE lab at the VUB.

MindXpres is based on a plug-in architecture. Plug-ins can do anything from arranging data in a certain way to letting the audience control the slideshow. Virtually anything is possible if you only implement it, and implementing it is fairly simple if you know a bit about web development as the whole thing is written in HTML5. Other software packages have plug-ins too of course, but they have a limited set of functionality available to them, they're not as easy to implement, and most importantly: they're bound by the same slide format used since overhead projections.

However, even with the new alternatives, Microsoft PowerPoint remains the most-used slideshow software. People keep using it because it's familiar, they've used it hundreds of times before and as such all their existing work is viewable only through PowerPoint. Switching to a new software package is hard. This thesis aims to make the transition easier, by providing a way to convert existing PowerPoint presentations into MindXpres. On top of that, we try to find a way to immediately release the transferred content from the confines of classic slides, by instead automatically figuring out the best possible layout for the content we extracted from the original PowerPoint file.

2.1 Terminology

The words *slideshow* and *presentation* are often used interchangeably throughout this report, although they do not quite cover the same meaning. By *slideshow* we mean a presentation consisting of a set of slides, the kind Microsoft PowerPoint and many other presentation software provide us with. *Presentation* then refers to the wider concept of material intended to be viewed and manipulated by people in order to convey information, usually but not necessarily from one or several presenter(s) to an audience.

The term *layout* refers to both the process of determining the position and size of each visual object that is to be displayed in a presentation, and the result of that process.

Slideware is a contraction of the words ‘slideshow’ and ‘software’, referring to software packages used to create slideshow presentations.

2.2 *Stuff we want to solve*

According to several sources (Parker, 2001, Drucker et al., 2006, Bajaj, 2013), over 30 million PowerPoint presentations are being made every day. That is an enormous amount. Creating a PowerPoint presentation is easy; creating a good PowerPoint presentation, however, is not. Slides have a fixed size, and you can only fit so much information on one slide before the effectiveness of transferring that information to one’s audience starts deteriorating. Over the years, many people have created written and unwritten guidelines to creating effective slideshows, specifying how much text and how many images should fit on one slide. Over those years, many people have failed to follow those guidelines. But whether you choose to follow the guidelines or not, one thing remains true: people who create slideshow presentations spend most of their time not on the *content* of their presentation, but on the *layout*.

The layout of a presentation can have a significant impact on how well it communicates information to and obtains information from those who interact with it. The vast majority of layouts created today are done “by hand”: a human graphic designer or “layout expert” makes most, if not all, of the decisions about the position and size of the objects to be presented. Designers typically spend years learning how to create effective layouts, and may take hours or days to create even a single screen of a presentation. Designing presentations by hand is too expensive and too slow to address situations in which time-critical information must be communicated.

Since layout is such a hard skill to master, we propose to automate this task, letting the presenter focus on the content of the presentation and providing a proper layout fit for the content provided through software.

2.2.1 *Real-life slideware problems*

It may seem like an overstatement to emphasize the significance of layout and formatting in presentations. One could assume these issues are irrelevant, or that only inexperienced presenters would make these mistakes. The real-life example of the space shuttle Columbia illustrates that this is not always the case. Leading up to the tragic incident in which the shuttle burned up during re-entry after spending 2 weeks in orbit, Boeing Corporation engineers

delivered three reports to NASA totalling 28 PowerPoint slides, to help them assess the damage caused by a piece of debris hitting the wing of the shuttle during launch, and the threat this damage might have posed. As Edward Tufte beautifully describes in his article “PowerPoint Does Rocket Science” (Tufte, 2005), the reports existed only in those slides, and the slides were woefully inadequate for the task at hand. Although Tufte likes to suggest this proves that PowerPoint is an inherently bad tool, what it really proves is that PowerPoint makes it easy to create bad presentations, and a tool that either discourages this manner of presenting information or makes it altogether impossible would be a great improvement.

3. RELATED WORK

This chapter’s content is largely based on “MindXpres: An Extensible Content-driven Cross-Media Presentation Platform” (Roels and Signer, 2014).

3.1 Background

The importance of digital presentations in this day and age cannot be understated. Millions of presentations are created every day, supporting the oral transfer of knowledge and playing an important role in educational settings. Their origins as tools for creating physical media such as photographic slides or transparencies for overhead projectors are still reflected in the underlying concepts and principles of slide-based presentation tools. The rectangular boundaries of a slide, and the linear navigation between slides, are still restrictions we face today in digital presentations. Tufte argues that these concepts of slideware have a negative impact on the effectiveness of knowledge transfer (Tufte, 2003). While the presenter is compelled to squeeze complex ideas into a linear sequence of slides, those ideas are rarely sequential by nature, resulting in a loss of relations, overview and details. An initial approach to address these issues might involve creating minimalistic presentations or introducing some structure via a table of contents. Sadly, when complex knowledge or other pieces of rich information need to be presented as is (Farkas, 2006) — as in the domain of learning — this does not work.

One of the main issues with traditional slideware presentations is their monolithic nature, especially when content is spread over many self-contained presentation files. “Reusing” previous work involves either switching between files while giving a presentation or duplicating some slides in the new presentation. It should be noted that this issue is not limited to the reuse of single slides: there is an ever increasing wealth of resources available for reuse, spread over a wide spectrum of distribution channels and formats. The possibility to include content by reference or transclusion (Nelson, 1995) may contribute in crossing the boundaries between different types of media and prove beneficial in the context of modern cross-media presentation tools.

The difference in functionality between the authoring of content and its visualisation is striking as well. The primary editors consist of mostly tool-

bars and buttons used for selecting and specifying the way content should be visualised, while support for authoring the content itself is not quite as extensive. Modern slideware has grown to include basic multimedia types such as videos, but most content is still rather static. It is, for example, not possible during a presentation to easily switch from a bar chart to a pie chart data visualisation, or to dynamically change some values in the represented data and immediately see the effect in the graph, which could be beneficial for knowledge transfer (Holzinger et al., 2008). The audience could also be more actively involved in the presentation, through audience response and classroom connectivity systems providing multi-device interfaces allowing to share knowledge and results during as well as after a presentation. The evolution of presentations is reminiscent of the Web2.0 movements where users have switched roles from purely consuming content to contributing as well, content has become more dynamic and interactive, and service-oriented architectures (“The Cloud”) have ensured decentralisation of content.

In order to move a step towards the next generation of cross-media presentation tools, it is essential to allow the rapid prototyping and evaluation of new concepts for the representation, visualisation and interaction with content.

After introducing existing slideware solutions, we discuss the requirements for next generation presentation tools. This is followed by a description of the extensible MindXpres architecture and its plug-in mechanism. The web technology-based implementation of MindXpres is validated based on a number of use cases and MindXpres plug-ins and followed by a discussion of future work.

A specific issue with slideware we’d like to focus on in this thesis, is the trouble with layout in presentations. It can be hard to display the content you want in a way that’s clear, informative and nice to look at. The vast majority of layouts created today is mostly done by hand: a human graphic designer or “layout expert” makes most, if not all, of the decisions about the position and size of the objects to be presented (Lok and Feiner, 2001). Most software offers some templates, allowing you to drop pictures and text into predefined slots and places on a slide, but then those templates have been defined by someone else too. Computer-generated layout is rare and usually not quite up to the task.

MindXpres is among the software packages offering templates, in that layout is handled by whichever plug-in you choose, but so far no plug-ins have defined dynamic layout algorithms, rather sticking to predefined ways to put text and pictures on slides. But as MindXpres does not constrain us to the limits of slides, this should be seen as an opportunity to offer dynamic layout as well. After all, if we’re not limited to a certain area within which

our content should fit, it should be much easier to put content next to each other in a way that makes sense.

3.2 *Existing solutions*

The impact, benefits and issues of slideware have been studied ever since digital slideware has been introduced. While some studies acknowledge slideware as a teaching aid (Holzinger et al., 2008), Tufte (Tufte, 2003) heavily criticises slideware for sticking to outdated concepts. He addresses the many consequences of spatial limitations or linear navigation and relates them to the human mind which works differently. One of Tufte's conclusions, which is also confirmed by Adams (Adams, 2006), is that slide-based presentations are not suitable for all kinds of knowledge transfer and in particular not in scientific settings. Recent work shows that it is important for the learning process that content is well integrated in the greater whole, both structurally and visually (Gross and Harmon, 2009), which is influenced by the navigation and visualisation. There have been a number of different approaches to offer non-linear navigation. Zoomable User Interfaces (ZUIs) as used by Counter-Point (Good and Bederson, 2002), Fly (Lichtschlag et al., 2009) or Prezi, offer virtually unlimited space. Also Microsoft has experimented with zoomable interfaces in pptPlex. While ZUIs are one way to escape the boundaries of the slide, we have seen other approaches such as MultiPresenter (Lanir et al., 2008) or tiling slideshows (Chen et al., 2006). PaperPoint (Signer and Norrie, 2007a) and Palette (Nelson et al., 1999) further enable the non-linear navigation of digital presentations based on a slide selection via augmented paper-based interfaces. Finally, there is a category of authoring tools that use hypermedia to enable different paths through a set of slides. NextSlidePlease (Spicer et al., 2012) allows users to create a weighted graph of slides and may suggest navigational paths based on the link weights and the remaining presentation time. Microsoft follows this trend with their HyperSlides (Edge et al., 2013) project. The potential of Microsoft PowerPoint as an authoring tool for hypermedia-based presentations has further been investigated by Garcia (Garcia, 2004).

Microsoft PowerPoint was officially released in 1990, with Windows 3.0 (Austin, 2009). It had originally been developed as Presenter, but trademark issues caused a name change early on. It was also originally build for the Macintosh, which may seem surprising nowadays but was actually common practice back then since the Macintosh was widely regarded as a better development environment, more mature, more stable and capable of far better performance and visualisations. Some may argue this still rings true today.

Since then, it has grown to be the world's most popular slide show presentation program, allegedly having been installed on over 1 billion computers worldwide, and being used on average 350 times *per second* (Parks, 2012). In 2012, it had a market share of 95%, leaving the other 5% to be shared by alternatives such as Apple's Keynote, Prezi, SlideRocket and others. While this number is declining, it may not be going as fast as many people think. As most readers of this thesis have heard before, over 30 million PowerPoint presentations are created every day, for all kinds of purposes, with good and bad results both presentation-wise and goal-wise.

Existing presentation tools require content to be duplicated for reuse, resulting in multiple redundant copies that need to be kept up to date. Even though some attempts have been made to address this issue, there is room for improvement. When it comes to document formats for more general educational purposes, there are formats such as the Learning Material Markup Language (LMML) (Sü and Freitag, 2002), the Connexions Markup Language (CNXML) and the eLesson Markup Language (eLML) (Fisler and Bleisch, 2006). The common factor of these formats is their focus on the reuse of content, but always at a relatively high granularity level. Content is organised in lessons or modules and users are encouraged to use these, as a whole, in their teaching. When we examined the formats in more detail, we noticed that they support outgoing links to external content but not the inclusion of content via references (transclusion). In the context of presentations, Microsoft's Slide Libraries are central repositories for the storage of slides in order to facilitate slide sharing and reuse within a company. However, one needs to set up a SharePoint server which might represent a hurdle for some users. Slides still need to be searched and manually copied into presentations. Furthermore, users are responsible to push back updates to the repository or update slides when they have been modified on the server side. Other commercial tools with similar intentions and functionality for content reuse are SlideRocket or SlideShare. The SliDL (Canós-Cerdá et al., 2010) research framework provides a service-oriented architecture for storing and tagging slides in a database for reuse. However, it shares some of the shortcomings of Microsoft's Slide Libraries. The ALOCOM (Verbert et al., 2008) framework for flexible content reuse consists of a content ontology and a (de)composition framework for legacy documents including PowerPoint documents, Wikipedia pages and SCORM content packages. While ALOCOM succeeds in the decomposition of legacy documents, it might be too rigid for evolving presentation formats and the tool is furthermore only supporting the authoring phase.

There is not only a similarity in the evolution of the Web and presentation environments, but a number of the issues presented in this section have

solutions in the setting of the Web. It is therefore not surprising that more recently we see the use of web technologies for realising presentation solutions. The Simple Standards-based Slide Show System (S5)¹ is an XHTML-based file format for slideshows which enforces the classical slideware model. The Slidy (Raggett, 2006) initiative by the W3C introduces another presentation format which is based on the standard slideware model. While these two formats are too limited for our needs, they have some interesting properties. Both formats show a clean separation of content and presentation via CSS themes. The visualisation is resolution independent and the layout and font size are adapted to the available screen real estate. Finally, we would like to mention recent HTML5-based presentation solutions, including projects such as impress.js, deck.js, Shower or reveal.js. A major benefit of applying a widely used open standard such as HTML is the cross-device support. Nevertheless, also these solutions show some restrictions in terms of visualisation, navigation and cross-media support.

Most of the tools and projects presented in this section focus on specific novel ideas for presentations. However, there is no interoperability between the concepts introduced in different tools. While one project might focus on the authoring, another one focusses on novel content types and a third solution introduces radically new navigation mechanisms. Some slideware tools can be extended via third-party plug-ins but the functionality that is exposed to the developers is often limited by the tool's underlying model. For example, PowerPoint allows plug-ins to interact with the presentation model, but the model dictates that a presentation must consist of a sequence of slides. This lack of freedom is also a shortcoming of existing web-based presentation formats. We therefore see a need for an open presentation platform such as MindXpres which supports innovation by providing the necessary modularity and interoperability (Bush and Mott, 2009).

It is perhaps surprising that, to our knowledge, currently no tools exist to calculate dynamic layouts of content in slideware. Existing solutions include template systems, sometimes very fine-grained like L^AT_EX allowing you to define templates for every single layout choice, usually more coarse like Microsoft PowerPoint or Apple Keynote using Master Slides to define different layouts on a per-slide basis, and always with the option of letting the user customize the layout by hand, literally manually moving the content to the exact place where we want it, unhindered by style guides, good practices or common sense. This has resulted in mindboggling layout choices involving enormous amounts of tiny text crammed onto one slide, or pictures strewn across a slide overwhelming the audience with too much information at once.

¹ <http://meyerweb.com/eric/tools/s5/>

3.3 New solutions

We now introduce a number of requirements to support a broad range of presentation styles and visualisations which have been compiled based on a review of the more recent presentation solutions presented in the previous section.

Non-linear Navigation As outlined earlier, the linear traversal of slides is a concept that has been taken over from early photographic slides. Nowadays, users are accustomed to this form of navigation even if it might come with some disadvantages. Any navigation outside of the predefined linear path (e.g. to answer a question from the audience) is rather complicated, since the presenter either needs substantial time to scroll forwards or backwards to the desired slide or has to switch to the slide sorter view. It is further impossible to include a single slide multiple times in the navigational path without any duplication. There are different ways how this lack of flexible navigation can be addressed, including presentation tools that allow the presenter to define non-linear navigation paths (Spicer et al., 2012) (Edge et al., 2013) or zoomable user interfaces (ZUIs) (Good and Bederson, 2002) (Lichtschlag et al., 2009) (Haller and Abecker, 2010).

Separation of Content and Presentation In order to facilitate experimentation with different visualisations, there should be a clear separation between content and presentation. This allows the authors of a presentation to focus on the content while the visualisation is handled by the presentation tool. Note that this approach is similar to the \LaTeX typesetting system where content is written in a standardised structured way and the visualisation is automatically handled by the typesetting system. There is also a \LaTeX document class for presentations called Beamer and we were inspired by its structured and content-driven approach. However, the content-related functionality and the visualisation are too limited to be considered as a basis for an extensible presentation tool.

Extensibility In order for a presentation tool to be successful as an experimental platform for new presentation concepts, it should be easy to rapidly prototype new content types and presentation formats as well as innovative navigation and visualisation techniques. It has to be possible to add or replace specific components without requiring changes in the core. In order to be truly extensible, a presentation tool should provide a modular architecture with loosely coupled components. Note that this type of extensibility should

not only be offered on the level of content types but also for the visualisation engine or content structures.

Cross-Media Content Reuse In the introduction we briefly mentioned the lack of content reuse in existing presentation tools. There is a wealth of open education material available but it is rather difficult to use this content in presentations. On the other hand, the concept of transclusion works well for digital documents and parts of the Web (e.g. via the HTML img tag). A modern presentation tool should also support the seamless integration of external cross-media content. This includes various mechanisms for including parts of other presentations (e.g. slides), transcluding content from third-party document formats as well as including content from open learning repositories.

Connectivity With the rise of social and mobile technologies, connectivity for multi-device input and output becomes more relevant in the context of presentation tools. Support for multi-directional connectivity is required for a number of reasons. First, it is necessary for the previously mentioned cross-media transclusion from external resources. Second, multi-directional connectivity forms the backbone for audience feedback via real-time response or voting systems (Dufresne et al., 1996) as well as other forms of multi-device interfaces.

Interactivity We mentioned that content might be more interactive and the extensibility requirement addresses this issue since the targeted architecture should support dynamic or interactive content and visualisations. Nevertheless, the use of mouse and keyboard might not be sufficient for components offering a high level of interaction. Therefore, a presentation tool should enable the integration of other forms of input such a gesture-based interaction based on Microsoft's Kinect controller or digital pen interaction (Signer and Norrie, 2010) as offered by the PaperPoint (Signer and Norrie, 2007a) presentation tool.

Post-Presentation Phase Even if it was never the original goal of slide decks, they often play an important role as study or reference material. While the sharing of traditional slide decks after a presentation is trivial, this changes when the previously mentioned requirements are taken into account. For instance, the nonlinear navigation allows presenters to go through their content in a non-obvious order or input from the audience might drive parts

of a presentation. Special attention should therefore be paid to the post-presentation phase. It should not only be easy to play back a presentation with the original navigational path, annotations and audience input, but its content should also be made discoverable and reusable. In accordance with the Web 2.0, we see potential for the social aspect in a post-presentation phase via a content discussion mechanism.

3.4 MindXpres Platform

In this section, we present the general architecture of our MindXpres² cross-media presentation platform which is outlined in Figure 3.1 and addresses the requirements presented in the previous section.




Fig. 3.1: MindXpres architecture

3.4.1 Document Format and Authoring Language

Content is stored, structured and referenced in a dedicated MindXpres document format. An individual MindXpres document contains the content itself and may also refer to some external content to be included. A new MindXpres document can be written manually similar to the \LaTeX approach introduced earlier or in the near future it can also be generated via a graphical authoring tool. In contrast to other presentation formats such as Slidy, S5 or OOXML, the authoring language eliminates unnecessary HTML and XML specifics and focusses on a semantically more meaningful vocabulary. The vocabulary of the authoring language is almost completely defined by plug-ins that provide support for various media types and structures. In order to give users some freedom in the way they present their information, the core MindXpres presentation engine only plays a supporting role for plug-ins and lets them define the media types (e.g. video or source code) as well as structures (e.g. slides or graph-based content layouts).

² <http://mindxpres.com>

This is also reflected in the document format as each plug-in extends the vocabulary that can be used. Any visual styling including different fonts, colours or backgrounds is achieved by applying specific themes to the underlying content.

3.4.2 Compiler

The compiler transforms a MindXpres document into a self-contained portable MindXpres presentation bundle. While a MindXpres document could be directly interpreted at visualisation time, for a number of reasons we decided to have this intermediary step. First, the compiler allows different types of presentations to be created from the same MindXpres document instance. This means that we can not only create dynamic and interactive presentations but also more static output formats such as PDF documents for printing. Similarly, we cannot always expect that there will be an Internet connection while giving a presentation. For this case, the compiler might create an offline version of a presentation with all necessary content pre-downloaded and included in the MindXpres presentation bundle. Last but not least, the compiler might resolve incompatibility issues by, for instance, converting unsupported video formats.

3.4.3 MindXpres Presentation Bundle

The dynamic MindXpres presentation bundle consists of the compiled content together with a portable cross-platform presentation runtime engine which allows more interactive and networked presentations. Similar to the original document, the compiled presentation content still consists of both, integrated content and references to external resources such as online content that will be retrieved when the presentation is visualised. Note that the content might have been modified by the compiler and, for example, been converted or extracted from other document formats that the runtime engine cannot process. References to external content may have been dereferenced by the compiler for offline viewing.

A presentation bundle's core runtime engine consists of the three modules shown in Figure 3.1. The *content engine* is responsible for processing the content and linking it to the corresponding visualisation plug-ins. The *graphics engine* abstracts all rendering-related functionality. For instance, certain presenters prefer a zoomable user interface in order to provide a better overview of their content (Reuss et al., 2008). This graphical functionality is also exposed to the plug-ins, which can make use of the provided abstractions. The *communication engine* exposes a communication API that can be used by

plug-ins. It provides some basic functionality for fetching external content but also offers the possibility to form networks between multiple MindXpres presentation instances as well as to connect to third-party hardware such as digital pens or clicker systems.

In addition to the presentation content and core modules, the presentation bundle contains a set of *themes* and *plug-ins* that are referenced by the content. Themes may contain visual styling on a global as well as on a plug-in level. When the content engine encounters different content types, they are handed over to the specific plug-in which uses the graphics engine to visualise the content.

3.4.4 Plug-in Types




Fig. 3.2: Structure plug-in examples

In order to provide the necessary flexibility, all non-core modules are implemented as plug-ins. Even the basic content types such as text, images or bullet lists have been realised via plug-ins and three major categories of plug-ins have to be distinguished:

- *Components* form the basic building blocks of a presentation. They are represented by plug-ins that handle the visualisation for specific content types such as text, images, bullet lists, graphs or videos. The content engine invokes the corresponding plug-ins in order to visualise the content.
- *Containers* are responsible for grouping and organising components of a specific type. An example of such a container is a slide with each slide containing different content but also some reoccurring elements. Every slide of a presentation may for example contain elements such as a title, a slide number and the author's name, which can be abstracted in a higher level container. Another example is an image container that visualises its content as a horizontally scrollable list of images. Note that we are not restricted to the slide format and content can be laid out in alternative ways.

- *Structures* are high-level structures and layouts for components and containers. For example, content can be scattered in a graph-like structure or it can be clearly grouped in sections like in a book. Both are radically different ways of visualising and navigating content but by abstracting them as plug-ins, the user can easily switch between different presentation styles as the ones shown in Figure 3.2. Structures differ from containers by the fact that they do not impose restrictions on the media types of their child elements and may also influence the default navigational path through the content.

3.4.5 Implementation

HTML5 and its related web technologies were chosen as the backbone for the MindXpres presentation platform. Other options such as JavaFX, Flash or game engines have also been investigated, but HTML5 seemed to be the best choice. The widely accepted HTML5 standard makes MindXpres presentations highly portable and runnable on any device with a recent web browser, including smartphones and tablets. Furthermore, HTML5 provides rich visualisation functionality out of the box and the combination with Cascading Style Sheets (CSS) and third-party JavaScript libraries forms a potent visualisation platform.

Document Format and Authoring Language

The MindXpres document format which allows us to easily express a presentation's content, structure and references is based on the eXtensible Markup Language (XML). A simple example of a presentation defined in our XML-based authoring language is shown in Listing 3.3. The set of valid tags and their structure, except the **presentation** root tag, is defined by the available plug-ins.

```
1 <presentation>
2   <slide title="Vannevar Bush">
3     <bulletlist>
4       <item>March 11, 1890 - June 28, 1974</item>
5       <item>American Engineer, founder of Raytheon</item>
6     </bulletlist>
7     <image source="bush.jpg"/>
8   </slide>
9 </presentation>
```

Lst. 3.3: Authoring a simple MindXpres presentation

Compiler

The compiler has been realised as a Node.js application. This not only allows the compiler to be used via a web interface or as a web service, but projects such as node-webkit also enable the compiler to run as a local offline desktop application. The choice of using server-side JavaScript was influenced by the fact that Node.js is capable of bridging web and desktop technologies. On the one hand, the framework makes it easy to interact with other web services and to work with HTML, JSON, XML and JavaScript visualisation libraries at compile time. On the other hand, the framework can also perform tasks which are usually not suited for web technologies, including video conversion, legacy document format access, file system access or TCP/IP connectivity.

In order to validate a MindXpres document in the XML format described above, there is an XML Schema which is augmented with additional constraints provided by the plug-ins. After validation, the document is parsed and discovered tags might trigger preprocessor actions by the plug-ins such as the extraction of data from referenced legacy document formats (e.g. PowerPoint or Excel) or the conversion of an unsupported video format. The tag is then converted to HTML5 by simply encoding the information in the attributes of a `div` element. The HTML5 standard allows custom attributes if they start with a `data-` prefix. Listing 3.4 shows parts of the transformed XML document shown in Listing 3.3. Note that the transformation does not include visualisation-specific information but merely results in a valid HTML5 document which is bundled into a self-contained package together with the presentation engine.

```
1 <div data-type="presentation">  
2   <div data-type="slide" data-title="Vannevar Bush">  
3     <div data-type="bulletlist">  
4       ...
```

Lst. 3.4: Transformed HTML5 presentation content

Presentation Engine

The presentation engine's task is to turn the compiled HTML content into a visually appealing and interactive presentation. As highlighted in Figure 3.1, the presentation engine consists of several smaller components which help plug-ins to implement powerful features with minimal effort. The combination of these components enables the rapid prototyping and evaluation of innovative visualisation ideas. A resulting MindXpres presentation combin-

ing various structure, container and component plug-ins is shown in Figure 3.5.



Fig. 3.5: A MindXpres presentation

Content Engine When a presentation is loaded, the content engine is the first component that is activated. It processes the content of the HTML presentation by making use of the well-known jQuery JavaScript library. Whenever a `div` element is discovered, the `data-type` attribute is read and the corresponding plug-ins are notified in order to visualise the content.

Graphics Engine The graphics engine provides support for interesting new visualisation and navigation styles. Next to some basic helper functions, it offers efficient panning, scaling and rotation via CCS3 transformations and supports zoomable user interfaces as well as the more traditional navigation approaches.

Communication Engine The communication engine implements abstractions that allow plug-ins to retrieve external content at run time. It further provides the architectural foundation to form networks between different MindXpres instances or to integrate third-party hardware (Roels et al., 2014). For our MindXpres prototype, we used a small Intel Next Unit of Computing Kit (NUC) with high-end WiFi and Bluetooth modules to act as a central access point and provide the underlying network support. MindXpres instances use WebSockets to communicate with other MindXpres instances via the access point. The access point further acts as a container for data adapters which translate input from third-party input and output devices into a generic representation that can be used by the MindXpres instances in

the network. In order to go beyond simple broadcast-based communication, we have implemented a routing mechanism based on the publish-subscribe pattern where plug-ins can subscribe to specific events or publish information. The communication engine provides the basis for audience response systems (Roels et al., 2014) or even full classroom communication systems where functionality is only limited by the creativity of plug-in developers.

Plug-ins Plug-ins are implemented as JavaScript bundles which consist of a folder containing JavaScript files and other resources such as CSS files, images or other JavaScript libraries. As a first convention, a plug-in should provide a manifest file with a predefined name. The manifest provides meta-data such as the plug-in name and version but also a list of tags to be used in a presentation. The plug-in claims unique ownership for these tags and is in charge for their visualisation if they are encountered by the content engine. As a second convention, a plug-in must provide at least one JavaScript file implementing certain methods, one of them being the `init()` method which is called when the plug-in is loaded by the presentation engine. It is up to the plug-in to load additional JavaScript or CSS via the provided dependency injection functionality. A second method to be implemented is the `visualise()` method which is invoked with a pointer to the corresponding DOM node as a parameter when the content engine encounters a tag to be visualised. A plug-in is free to modify the DOM tree and may also register callbacks to handle future interaction with the content.

Themes We currently use CSS to provide a basic templating system. These themes offer styling either on a global or on a plug-in level. However, we see this as a temporary solution as it is not well-suited for alternative compiler outputs (e.g. PDF) and a more generic templating scheme is planned for the future.

3.4.6 Use Cases

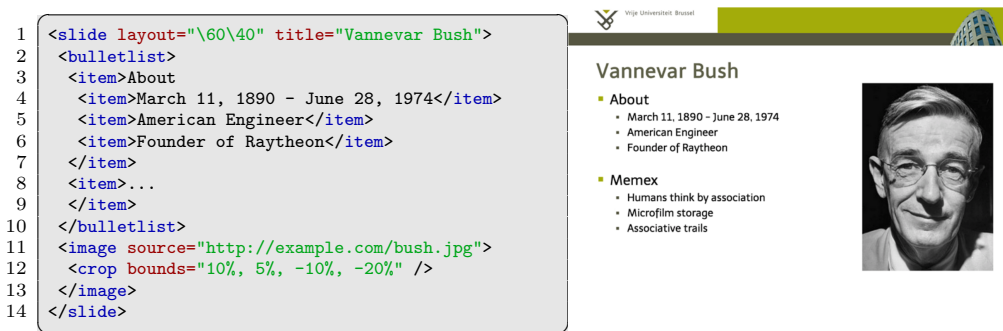
In order to validate the architectural and technological choices, we demonstrate the extensibility and feasibility of MindXpres as a rapid prototyping platform by presenting a number of content- and navigation-specific plug-ins that have been developed so far. Additional plug-ins for audience-driven functionality such as real-time polls, screen mirroring and navigational takeover can be found in (Roels et al., 2014).

Structured Overview Plug-in

In Section 3.4 we have explained how structure plug-ins may change the way presentations are visualised and navigated. In order to illustrate this, we have implemented a structure plug-in called *structured layout* which combines a zoomable user interface with the ability to group content into sections. The resulting visualisation of the *structured layout* plug-in is shown in Figure 3.5. Whenever a new section is reached, the view is zoomed out to provide an overview of the content within the section and communicate a sense of progress.

Slide Plug-in

In order to also support the traditional slide concept, we created a slide-like container plug-in. While the benefits and issues of using slides with a fixed size are debatable, we implemented this plug-in as a proof of the framework's versatility. The main function of the slide plug-in is to provide a rectangular styleable component container with a title and some other information. Containers may also offer functionality to layout their content. In this case, the slide plug-in offers a quick and easy layout mechanism which allows the presenter to partition the slide into rows and columns. Content is then assigned to these slots in the order that it is discovered. The use of the slide plug-in together with the resulting visualisation is exemplified in Listing 3.6. It also shows the use of the image plug-in (a component plug-in) which enables a simple form of cross-media transclusion. A visualised external image can be cropped and filters (e.g. colour correction) may be applied without duplicating or modifying the original source.




Lst. 3.6: Slide plug-in

Enhanced Video Plug-in

When videos are used in educational settings, we often need more functionality than what is offered by the average video player (Reuss et al., 2008). MindXpres provides the enhanced video plug-in shown in Listing 3.7 with the possibility to overlay a video with text or arbitrary shapes. This overlay functionality can be used as a basic captioning system as well as to highlight items of interest during playback.

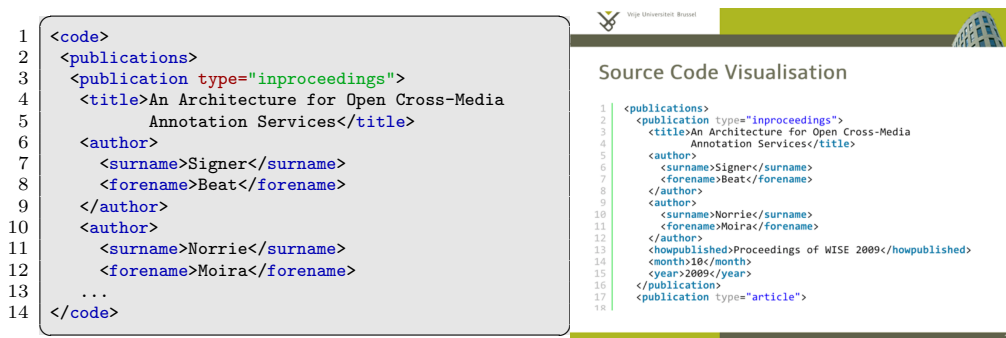
Furthermore, we added the option to trigger certain events at specified times. One can define that a video should automatically pause at a certain point, highlight an object and continue playing after a specified amount of time. Additional features include the bookmarking of certain positions in a video for direct access or the possibility to display multiple videos in a synchronised manner. Our enhanced video plug-in injects the default HTML5 video player and overlays it with a transparent `div` element for augmentation. Currently we make use of the HTML5 video API to synchronise the creation and removal of overlays but a SMIL-based implementation might be used in the future.

<pre> 1 <video source="vid.mp4"> 2 <caption start="0:00" duration="1500ms"> 3 Lecture 3 - Butterfly Species 4 </caption> 5 <pause start="0:43" duration="5s"> 6 <caption> 7 The peacock butterfly (aglais io) ... 8 </caption> 9 <highlight x="30%" y="9%" 10 width="35%" height="40%" /> 11 </pause> 12 </video> </pre>	
--	--

Lst. 3.7: Enhanced video plug-in

Source Code Visualisation Plug-in

Earlier, we mentioned the difficulty of visualising complex resources such as source code. Our MindXpres source code plug-in exports a `code` tag which allows the presenter to paste their code into a presentation and have MindXpres visualise it nicely by making use of syntax highlighting via the SyntaxHighlighter³ JavaScript library. Whenever the content engine encounters a `code` tag, it invokes the code plug-in to beautify the code and automatically adds vertical scrollbars for larger pieces of source code as shown in Listing 3.8.



Lst. 3.8: Source code visualisation

3.4.7 Discussion and Future Work

MindXpres currently supports transclusion and cross-media content reuse on the plug-in level. For instance, the image or video plug-in can visualise (and enhance) external resources, a dictionary plug-in might retrieve definitions on demand via a web service or we might create a plug-in that allows us to import content (e.g. PowerPoint slides) from legacy documents at compile time. Nevertheless, we are currently investigating the introduction of generic reuse tags in our document format which would allow the presenter to transclude arbitrary parts of other MindXpres presentations. While our focus has been on the cross-media aspect of resources that can be used in a presentation, we might also investigate the cross-media publishing aspect via alternative compiler output formats.

We are aware that the current authoring of MindXpres presentations has some usability issues. The average presenter cannot be expected to construct an XML document or any CSS themes. In order to tackle this issue and further evaluate MindXpres in real-life settings, we are currently developing

³ <http://alexgorbatchev.com/SyntaxHighlighter/>

a graphical MindXpres authoring tool. We further intend to provide a central plug-in repository which would make it easy for novice users to find, install and use new plug-ins via the authoring tool. In the long run, we intend to revise the use of monolithic documents and move towards repositories of semantically linked information based on the RSL hypermedia metamodel (Signer and Norrie, 2007b). This would not only promote content reuse and sharing, but also create opportunities for context-aware as well as semi-automated presentation authoring where relevant content is recommended by the authoring tool.

3.4.8 *Conclusion*

We have presented MindXpres, an HTML5 and web technology-based presentation platform addressing the lacking extensibility of existing slideware solutions. The extensibility of the MindXpres platform heavily relies on a plug-in mechanism and we have outlined how the tool can be extended on the content, visualisation as well as on the interaction level. By providing different forms of cross-media content transclusion, our solution further avoids the redundant storage and replication of slides. The presented MindXpres document format in combination with a compiler and the corresponding plug-ins offers the opportunity to have compile- or run-time cross-media content transclusion from third-party resources. At the same time, the flexible and extensible document model in combination with a zoomable user interface allows us to escape the boundaries of traditional slide formats. The presented MindXpres solution represents a promising platform for the unification of existing presentation concepts as well as for the rapid prototyping and investigation of new ideas for next generation cross-media presentation solutions.

MindXpres tries to be an alternative unlike other alternatives, stepping away from the classic slide format and introducing a plug-in system to allow literally anything you can think of. The obvious choices are visualisations, animations, layout and the embedding of various media such as video, but it can also provide interaction with the audience, allow different presenters to take control from their own device, arbitrarily and ad-hoc change the way data is presented allowing to play into the audience's reactions unhindered by design decisions made beforehand, and so many other novel ways of presenting that captivate the audience instead of lulling them to sleep.

MindXpres was first created by my advisor, Reinout Roels, in 2012 as part of his Master thesis, and has since been built upon by himself and other students.

3.5 Layout

Proper layout is incredibly important when trying to transfer knowledge and information through written and visual media. Layout can help clarify boundaries and relations between pieces of information, by grouping and separating them appropriately. Layout is one component of a presentation's design, that — combined with other decisions — determines the number and nature of the visual representations of the information the creator wants to communicate, along with its format⁴. The layout of a presentation can have a tremendous influence on its effectiveness in communicating information to, and obtaining information from, the audience it is meant to interact with. The importance of individual objects can be emphasized or minimized, and the connection between objects can be clarified or blurred. A well laid out presentation can provide a narrative for the viewer to discover, inferring correct links between the objects along the way, and to accomplish tasks quickly and correctly, increasing the presentation's effectiveness.

Creating a good layout is almost never easy. People often spend more time on the layout of their presentation than the content. Most, if not all, decisions in layout are made by human beings. Some of them are professional designers who spend years learning and figuring out how to create effective layouts, and even then they may take hours or days to create even a single screen of a presentation. In fact, the more someone knows about proper layout and design, the more time they may spend perfecting their work. However, sometimes time-critical information must be communicated and the layout process is too expensive and too slow to address these situations. This can be a serious problem (see also section 2.2.1). Many software packages have been developed to make this process easier, to get better results, to give more or less control to the creators. Many different approaches have been taken, and yet most of them still involve having a human being make the final decisions on the layout.

There are tools like PowerPoint, which give you some guidelines and some templates but generally let you do your own thing. If your own thing is entirely different from any best practices on layout, nothing will stop you. Other tools like L^AT_EX give you complete control over every aspect of the layout, while setting some sensible defaults so that you can get a good-looking layout without much effort, while still letting you do whatever you want once you overcome the steep learning curve that separates the casual users from the experts. There are tools that combine the power of L^AT_EX

⁴ The way the visual objects are realized (e.g. as text, graphics, UI widgets...), and their attributes (e.g. color, texture, font...)

with the comfort of WYSIWYG editors, bringing the casuals a bit closer to the experts. But all of those tools have one thing in common: every aspect of every layout they create has, at some point, been designed and decided upon by a human being.

Aesthetics are a natural phenomenon, and the creation of aesthetically pleasing layouts is therefore a manifestation of our instincts. As with most instincts, it has proven difficult to translate this into a concept that can be understood by a computer. Moreover, it is still difficult to explain it in human terms, which — according to a popular quote often attributed to Albert Einstein — proves we don't fully understand it ourselves.

When we look to other technologies, we do find some automated layout implementations. For example, the web has had to adapt to mobile devices with small screens over the past few years, and has done this gracefully by creating the concept of responsive design. In short, this allows websites to adapt their layout to any screen, no matter the size. While this is often a hard-coded difference, where effectively two or more versions of the same webpage are created aimed at different screen sizes, some websites take a more dynamic approach based on constraints. As the space the page is to be displayed on gets smaller, the layout algorithm may decide to display content below other content instead of side-by-side, it may scale images to fit the screen, it may even switch fonts and font sizes if necessary.

This constraint-based technique is described in a few papers (Lok and Feiner, 2001, Hurst et al., 2009), but has — to our knowledge — not been applied in any presentation software so far. This is surprising, as presentations often look like they could use some of this magic. A proper constraint-based layout algorithm could allow any user to drop content onto a slide, without worrying about clarity or even legibility, and the algorithm could take care of the rest. Of course, there are some limits in traditional slideware that may hinder this approach: if a user decides to put more content on a slide than there is physical room available, the algorithm could either make the content smaller or split it across several slides, but either solution may bring its own problems up. An advantage of ZUI's is that no matter how small the content gets, we can still zoom in to make it clear again⁵.

⁵ It should be noted that MindXpres in its current form does not support this level of zooming. While the software can zoom out to provide an overview of the presentation while zooming in on the separate components, it is not yet possible to zoom in or out extremely to reveal 'hidden' parts of the presentation. This is something we encourage to look into and change, because it can greatly improve both our layout solution as well as the whole MindXpres experience in general.

3.5.1 *Algorithms*

When researching layout algorithms, one will often come across the very active field of graph layout (Di Battista et al., 1998). We will not go into the specifics of this field, as most of the issues with which it is concerned are specific to problems caused by the explicit visual representation of graph edges — for example, the minimization of edge crossing (Di Battista et al., 1990, Shahrokhi et al., 1996). The same applies to automated layout as referring to automated circuit layout for VLSI chip fabrication (Hu and Kuh, 1985, Lengauer, 1990) as well as automated placement of pieces to be cut from a bolt of cloth used to produce clothing (Milenkovic et al., 1991). Contrary to presentation layouts (including graph layouts), these layouts are designed to meet the requirements of a fabrication process, rather than to make them understandable to humans. While some techniques used therein definitely apply to our more general problem of automated presentation layout (e.g. general constraint solvers) others decidedly do not (e.g. bin-packing techniques (Hofri, 1980) that result in minimal area layouts at the expense of maintaining visually obvious relationships between objects).

4. APPROACH

In this chapter we explain the different approaches we tried in order to reach our goal and find a solution for the problem we described. As you will see, this was not immediately a straightforward process but rather one of trial and error. The goal was clear, the starting point was clear as well, but as often in computer science, there is more than one way to get from point A to point B, and it is not always clear which way is the best, easiest, most efficient or most effective.

Since we're talking about the approach here, and not the implementation (for that, see chapter 5), we start by describing in broad terms what needs to be done and how this should be done, then we refine until we have a full set of specifications ready for implementation, where the last details will be ironed out.

Unfortunately it is possible to refine an approach until it is ready for implementation, and only find out during implementation that the approach you've chosen will not work. This happened during our work on creating an automated layout system. Luckily we still had time to go back to the drawing board, and we did not have to restart from scratch; large parts of our approach were correct, the basic layout process we thought out was still a viable part of the approach, but it turned out we would have to split up the conversion and layout parts into two separate processes, rather than implementing them as two steps of the same process.

Specifically, we had thought at first to figure out the ideal layout during conversion, when we would have all the separate components, by immediately putting them in the right place. This idea was partly conceived after looking at the HTML code generated by the MindXpres compiler, thinking we would generate the same HTML code in our conversion process. It turned out we could bypass the MindXpres compiler this way, but that wouldn't be necessary: we could just as well generate MindXpres XML and have the compiler take care of the rest for us.

We also found during implementation that generating a layout in Java would not easily give us the results we were hoping for. However, at this point we had realized generating MindXpres XML would be a better option, so we could have MindXpres take care of the layout for us. Except MindXpres

didn't do fully automated layout yet, the layout system was mostly template-based, so we decided we would need to write our own MindXpres plug-in that would solve this problem for us.

4.1 Conversion process

The first part of the approach is fairly straightforward in its basic explanation: we had to convert PowerPoint presentations into MindXpres presentations. This involved finding out how PowerPoint presentations are structured, getting the parts we need out of that structure, and then putting those parts together in the MindXpres structure.

It appeared soon enough to us that the nature of this process resembled that of a compilation process. A compiler takes source code and transforms it into a working program with the semantics described by that source code. The compilation process consists of several steps. First the source code is tokenized, which means the symbols in the code are identified one by one and classified in certain categories.

Then the tokens are processed by a parser into an intermediary form called a parse tree. A parser looks for certain predefined patterns in the source code. These patterns are part of the source code's language syntax. As such, these two steps analyse and validate the source code's syntax. If part of the code does not match any pattern, the parser and the compilation process stop and the user gets a message saying the code's syntax is invalid.

When a parse tree is constructed, the compilation process can alter it, to improve it. Certain patterns in the parse tree may be replaceable by different patterns with the same outcome, but with more optimal execution. This part of the compilation process is optional, and is called compiler optimization. Optimizations can consist of many things, depending on the language. For example, some languages guarantee tail call optimization, where infinite loops can be constructed by letting a function call itself as its last statement without causing a stack overflow. This is something the compiler (or interpreter) can optimize during this part of the compilation process.

After this, the parse tree can be written out to produce the desired output. Every node in the tree has a well-defined equivalent in the target language's syntax. The target language can be Assembly, which consists of the exact instructions a CPU needs to carry out a program, or it can be another programming language. Many compilers of higher-level languages translate their language into C, for several reasons: the C compilers that translate C into Assembly have been optimized so much that it is easier to rely on them than to put an enormous amount of effort into optimizing another language;

C compilers exist for most — if not all — CPU architectures, which means translating a language into C makes it compatible with all those architectures, while it would cost a lot more effort to write different compilers for every architecture you would want to make your language available on.

The conversion tool that is the purpose of this thesis, can be described in a similar succession of steps. As a first step, we take a PowerPoint presentation and take it apart into its components, effectively walking over each component, classifying them and registering their content type, original position and size, and any other specific properties. This can be seen as the tokenization phase, after which we end up with a series of ‘tokens’ or, in our case, presentation components.

We then turn this series of ‘tokens’ into a ‘parse tree’, an intermediary structure that reflects the relation between the components and the hierarchy of the presentation, which may consist of chapters, sections, slides and component groups. In PowerPoint this structure is fairly simple, so the creation of this ‘parse tree’ is a straightforward process.

However, in MindXpres we are not limited to the rigid hierarchy of sections and slides, so at this point we can actually start manipulating our tree and improve upon it, for example by moving parts around, nesting components in different ways, grouping them in other ways than they originally were, etc. In compilation terms, this is the optimization phase, where the compiler can manipulate the program to run more efficiently, to replace parts of it with other functionality, or to add features the source didn’t explicitly specify (e.g. garbage collection, but also spyware components (Scahill and Begley, 2015)).

As we discuss in section 4.2, this seemed like the right time to incorporate automated layout generation into the conversion process. As we see later in section 4.3.1, it turned out it wasn’t. In the end, no significant ‘optimizations’ or manipulation of the tree structure were implemented. Later on we would utilize this optimization phase to enable automated layout in another way, without actually performing the layout here, but at this point it would not affect the end result in any way.

To finish the conversion process, we can traverse our component tree and generate a MindXpres presentation from it. This can be done in several ways, since our intermediary form is in no way dependant on or bound to a specific format. Since the MindXpres compiler was unavailable for a long time during our research and implementation, we decided it would be best to go straight to HTML5, so that we could test the conversion process without relying on the MindXpres compiler. This worked out fairly well, although manually constructing HTML5 to work with the MindXpres JavaScript library proved difficult. We ran into several issues, often mostly due to our lack of knowledge

of the inner workings of MindXpres, but we managed to get a presentable result that emulated the original PowerPoint presentation quite well.

Afterwards, we altered our conversion tool to generate MindXpres XML instead, which was a lot simpler since we would rely on MindXpres to provide our layout and other things for us through the MindXpres compiler. This approach allowed us to use the full power of MindXpres, including our own plug-in for automated layout. At this point, the optimization phase was also revisited, and leveraged to introduce specific XML tags around component groups that would trigger our automated layout plug-in.

4.2 Compiler optimizations

Since the conversion process resembles that of a compiler, it seemed logical at first to make automated layout a part of that process, as some kind of ‘compiler optimization’. During this phase in the process, the component tree would be manipulated and altered, with the express purpose to improve upon its structure and properties, so as to get a better end result. Our improvements in this case would then consist of the automated layout.

As a first attempt, we tried to traverse the component tree, giving each object new coordinates and sizes based on their original coordinates and sizes, as well as the coordinates and sizes of objects around them, so that they would fit together on every slide as well as possible. This seemed an easy solution, but the results were sub-optimal. On top of that, we soon realised that we were in essence creating another template-based system that would generate slides and presentations based on predefined ratios and rules, which was exactly the opposite of what we were trying to do. As such, we abandoned this approach in favor of a constraint-based algorithm as described in section 3.5.1.

This involved a technique that at first sight may seem like yet another template system, but actually is completely different: defining constraints for every component, in the form of margins, maximum sizes and other limits, and then calculating a way to satisfy all constraints while fitting content together on each slide. The similarities with template-based systems exist in the presence of predefined constraints, ratios and rules, but the important difference is that these constraints are defined relative to the component itself, without specifying anything absolute about location or size. For example, we would retain the aspect ratio of an image, without specifying its size, so that the image may be scaled to accomodate other components in a dynamic layout. As another example, we might specify there needs to be a certain distance between a component and any other components, relative

to its size. We could also specify a certain relation between components, ensuring components stay in each others vicinity, one should always be left of the other, no other components may be placed between them, etc. Using these rules, we would then programmatically calculate the best layout using those components, but without any other bias. These constraints would be based only on the original situation, never on any suggestions from us or other developers or authors, which makes all the difference with traditional template-based layouts.

While this is clearly a better method, it turned out the compiler optimization phase was not the best place in the process to take care of this. While we had the necessary data to calculate the layout, we would have had to generate the layout along with the MindXpres presentation, after which the presentation could not be altered anymore without breaking the layout. This defeated the purpose of exporting to MindXpres, which was to allow the presenter to edit, extend and improve their presentation further using MindXpres. What we needed was a way to get MindXpres itself to generate the layout, even if we wanted to add components to the presentation afterwards, and even if we wanted to create a new MindXpres presentation instead of starting from PowerPoint. After all, how would we convince people to drop PowerPoint for MindXpres's automated layout capabilities if they could only use that functionality by starting from PowerPoint?

In the end, we decided to change our approach again. We took the automated layout out of the conversion process, instead opting at this point in the process to only add the necessary layout triggers in the form of an enclosing XML tag around the components that would need to be included in the automated layout. As such, the generated MindXpres XML would include those tags, and a plug-in (described in section 4.3.1) would then generate the layout at runtime.

4.3 Using MindXpres

One of the primary goals of MindXpres is to separate content from layout, allowing the author of a presentation to focus on the content while MindXpres takes care of the layout. The way it does this is currently mostly through the compiler, which decides the width, height and coordinates of content, relative to the container the content belongs to. The plug-ins responsible for handling components and containers currently don't mess with those settings, but technically, they could. The compiler decides the measurements and coordinates based on templates. The solution we were looking for was a layout engine that could take any content and put it in an appropriate

layout without any directions from the user. As such, we had to enhance MindXpres's layout engine to use constraints, based on the size of the content, and try to find an optimal position for every component it is given.

4.3.1 A MindXpres plug-in

We did this by creating an invisible container plug-in. Containers are a way of grouping components, other containers, etc. in MindXpres. This means they have control over their child elements, which gives us the opportunity to override the layout of those elements. A container plug-in thus allows us to implement our own layout system. Since it's a new element, it doesn't override existing elements as it would have done if we had, for example, rewritten the 'slide' plug-in. The user can decide for themselves whether or not to use it, and it can be used anywhere in the presentation: wrap the whole presentation in it, or just a small part, whichever works best for your purposes. It also won't break existing presentations that don't use it, while those presentations can very easily be altered to take advantage of it.

An important aspect of this is that containers can be nested. This means we can create slide-based presentations, which can contain our `autolayout` container, which then contains the slide's contents, thus creating an optimal layout of the content per-slide. Another way of using it could be without slides, throwing all content together in one `autolayout` container, and letting it take care of the layout for the whole presentation at once. It should be noted here that the `autolayout` container makes each of its child nodes focusable separately, to compensate for arbitrary resizing it may perform on large objects in order to fit them next to other content, by using the focus functionality to automatically zoom into these components when necessary.

We call it an *invisible* container plug-in because it does not introduce any visual content, shape or indication for itself. Compare with the `slide` plug-in which obviously puts some kind of slide-look around the content it encompasses, and it becomes clear what we mean by this: although the content within is obviously affected by our plug-in, there is no visible indication of its presence to the audience.

The plug-in uses the compiler's numbers to decide relative locations between components, as well as size ratios, and then finds a way to display those components in a way that the display order makes sense (or at least matches the intended order as closely as possible), that no overlapping occurs (since we don't have the animations that PowerPoint might have used to display one piece of information and then another on top of it), and resizing everything if necessary in order to fit within the specified container. While this may seem like a bad idea since content can get illegibly small this

way, keep in mind that we can rely on the ZUI¹ to focus on each component separately, or on groups of components, while PowerPoint obviously can only display the whole slide at once.

In this manner we would generate MindXpres presentations that were immediately usable, while also being adjustable; and on top of this, we would allow the automated layout process to be used in other MindXpres presentations that were not originally converted from PowerPoint slides. The goal of this plug-in would thus be to provide automated layout functionality to MindXpres presentations, and to allow any MindXpres author to use it simply by putting an `<autolayout></autolayout>` container around the components they want the plug-in to act on. This approach has the additional advantage that the container can be used multiple times throughout the presentation, while also allowing other parts of the presentation to have a manual layout.

Since it is possible to nest containers, which means a number of components could be grouped together in an `autolayout` container, then the result could be put into another `autolayout` container together with other components — other `autolayout` containers, perhaps — to generate an automated layout for an overview of the different groups. Compare it with traditional slideware, where components are grouped together in slides, then the slides might be put next to each other in an overview — except our approach drops the slide boundaries, while still maintaining the ability to group components together to show a relation or link between them.

4.3.2 An automated layout algorithm

As discussed earlier, our first approach included an algorithm where content would be placed on slides according to certain rules, trying to attain a mythical ‘perfect’ layout based on the golden ratio, symmetry, centering and other general guidelines we would find in advice on creating presentations. It turns out that, while following those guidelines as a human being is generally a good idea, a computer has different ways of calculating a good layout. The issue here can be compared to other problems in computer science; for example, people in the robotics department have tried for decades to create a robot that behaves exactly like a human being, and people in artificial intelligence have tried to create an AI that thinks like us. However, we’ve found time and time again that computers simply aren’t very good at acting, thinking or being human, just like we aren’t good at being computers. Making a computer act like a human makes it disadvantaged — almost by definition, just

¹ Zoomable User Interface

like we are severely handicapped when we try to perform typically automatable, repetitive and/or math-intensive tasks. A computer's true power only shows when you let it do what it's good at, which is the repetitive stuff, the mathematically complex stuff, etc. Trying to make it generate presentation layouts like a human would, is asking for subpar results.

If we approach this problem keeping in mind a computer's strength and weaknesses, we arrive at a different approach. This involves calculating sizes, ratio's, positions, margins and other numbers, of which the formulas are actually not too hard to come up with as a human, but which the execution is definitely more of a computer task. We start off by checking each component, and noting its original location and size. We then try to find components that are in proximity of each other, and figure out their original layout: above/below each other, next to each other, overlapping... Then, we try to put them together, possibly resizing them to match each other's sizes, and trying to match their original relative locations while introducing a certain rigidity, or consistency, by aligning them properly and puzzling them together as neatly as possible.

This last part may sound weird, but it really is something to take into consideration, especially when the amount of components might be much bigger than what should fit on an average traditional slide. You could put all components in a row, just displaying them side-by-side, but that is not very aesthetically pleasing. Instead, we opted to try and keep components close to each other. This was finally achieved by finding the location closest to the starting point that would fit the component being considered, while still taking into account the earlier constraints about relative location and size. Thanks to the ZUI in MindXpres, this makes for interesting layouts that still remain manageable, and provide a nice overview of all content when zoomed out.

5. IMPLEMENTATION

To implement the *ppt2mxp* conversion tool that is the subject of this thesis, we chose the Java programming language (Gosling and McGilton, 1995), version 8. Although the author has significant experience with lots of other, more interesting, more compelling, more fun languages, several reasons pushed us towards Java, the least of them being its ease of use. Of course, Java *is* easy to use — it would not have become as popular as it is nowadays if it wasn't. It has a fairly clear and logical syntax, a consistent structure, and an extensive standard library. At conception in 1995, its performance was abysmal, but throughout the years it has steadily improved and somewhere between Java 5 (then still called 1.5) and 6 (when they dropped the '1.' prefix) it became an industry standard.

Quite a number of IDEs have been created to further improve developers' experience working with Java. Netbeans, Eclipse and IntelliJ come to mind, although there are many others, and of course you can still write Java using a standard (or advanced) text editor such as Notepad or VIM. While the author usually prefers the latter for any kind of text editing — this very document was written entirely using VIM — the weapon of choice when it comes to Java is currently IntelliJ. The way IntelliJ practically writes more than half of the code automatically for you is something no other IDE has been able to match. Naturally, this is the author's personal opinion and should not be seen as fact, but if you're looking for a new Java IDE, it's definitely worth checking out. The prospect of using IntelliJ for this thesis has definitely contributed to the decision of using Java. It should be noted that, had another Java IDE been required, this thesis might never have seen the light of day.

The vast and extensive amount of libraries available for Java was obviously one of the more important reasons to make this choice. The existence of the Apache POI library (see section 5.1) was a huge help in reaching our goal; without it, we would have had to figure out the very obfuscated .ppt file format structure, which undoubtedly would have taken up more time than was available to us. Other libraries like Spring, which allows the programmer to use and reuse components without writing complex systems to instantiate them, further increased our resolve to make Java our primary technology

choice.

However, Java is not the only technology used here. MindXpres is written entirely in HTML5, so any tool that somehow relates to MindXpres sooner or later needs to use HTML5 as well. The widely accepted HTML5 standard makes MindXpres presentations highly portable and runnable on any device with a recent web browser, including smartphones and tablets (Roels and Signer, 2014).

In the following sections we discuss how the various technologies were used to create the *ppt2mxp* tool.

5.1 Taking PowerPoint apart

When converting one file format into another, the first part of the process involves getting the data you need out of the original file. This can be very complicated, as some — usually proprietary — file formats are deliberately designed to discourage this. They obfuscate data, encrypt it, and structure it in illogical and unexpected ways, amongst other techniques. The PowerPoint file format unfortunately is such a format, as Microsoft wouldn't want to risk other companies making software that would work with PowerPoint files. Of course, over the years people have managed to crack the format, enabling the conversion of PowerPoint presentations into other formats, although the conversion does not usually guarantee to yield results that mimic the original version perfectly. Luckily, we don't want a perfect conversion, we want a better one.

We found Apache POI library very helpful in this part of the implementation. The POI¹ Library is a Java library that provides an API to access Microsoft document formats. The most mature (and most popular) part of it is HSSF², which is used by Java developers worldwide to access Microsoft Excel spreadsheet data, as well as export data into Excel spreadsheets.

For our purposes, we relied on HSLF³, which provided us with a full API to access the contents of a PowerPoint presentation's contents in a myriad of ways. We could access all images at once, or every bit of text from the whole presentation, but the most interesting to us was the ability to access contents on a per-slide basis. Getting a list of the slides in a presentation first allowed us to group contents within their immediate context, under a node per slide in our component tree. As such, we could loop over the presentation's slides, converting them one by one, by placing the contents of

¹ Originally “Poor Obfuscation Implementation” (Sundaram, 2004)

² Horrible SpreadSheet Format

³ “Horrible SLideshow Format”

each slide in a MindXpres slide equivalent.

5.1.1 Bullets

That was unfortunately not the end of it. While HSLF does give us access to all the text in a presentation, or per slide, it was not immediately clear to us how it distinguished between ‘normal’ text and bullet lists. This meant for a long time our conversion process was incomplete, as all bullets from the original PowerPoint presentation appeared as incoherent text runs in our converted result. We found out about the `RichTextRun` class, which had all the tools and properties to detect bullets and their indentation level, but we only discovered very recently that we could extract `RichTextRuns` from the `TextShapes` we were getting out of the slides.

5.1.2 Animations

Another challenge was dealing with animations and other ways people managed to put way more content on one slide than would be advisable. The animations could not be transferred to MindXpres since MindXpres has its own way of transitioning from each component to the next in the form of a ZUI⁴. It would technically be possible to implement additional animations as a separate plug-in for MindXpres to provide the equivalents of the animations in Microsoft PowerPoint, but that is beyond the scope of this thesis. So we could not provide the same animations, but some people use those animations not just to show off but to actually show multiple pictures and blocks of text, one after the other, on the same slide. Without animations, this content would either not be visible or it would become a serious layout issue in MindXpres.

Our first solution tried to limit the amount of objects one slide can contain, and any additional content should be put on extra slides automatically. A downside of this is that we had no way of guessing the correct order in which the content should appear, so what may have been an intrinsic choreography of pictures in PowerPoint might become an incoherent jumble of images in MindXpres. Another solution would be to scale all content until it all fits next to each other on one slide, and then rely on the ZUI to show the pictures one by one, but in this case the same problem with order of appearance manifests itself. In the end, we decided it would be best to accept that no conversion algorithm is going to be perfect, and the author can always manually change the order around after the conversion is done.

⁴ Zoomable User Interface

With this in mind, we now render the components in the order we get them from HSLF, hoping that this resembles the original order closely. The automated layout takes care of any overlapping that might have occurred originally, so we don't have to worry about that.

5.2 Generating MindXpres

Generating MindXpres presentations was the final goal of the first phase of this thesis. This seemed a fairly easy task at first, until we learned that the MindXpres compiler would not be available to us for most of the year. This meant we would either not be able to view-test our generated presentations, or we would have to convert them to browser-ready HTML5 ourselves. We chose the latter option, as not being able to see our results would not be very helpful in implementing and tweaking our conversion tool. As a result, this task became much more complicated, as we had to emulate the compiler's work ourselves. Luckily we already decided we would be working with a Java object representation of the original presentation as an intermediary form, a so-called component tree, which meant we could easily change the output of our conversion tool without affecting the rest of the conversion process, and on a per-component basis.

5.2.1 Playing MindXpres compiler

Since the MindXpres compiler was not functional during most of this thesis' implementation, we decided to generate an HTML5 file much like the MindXpres compiler would, including the MindXpres JavaScript library and plug-ins. This required us to first learn how MindXpres works on the inside, which proved to be a steep learning curve but gave us more insight into the software than we would've gotten if we only had to generate MindXpres XML and leave the rest to the compiler.

Improving the ZUI

As an exercise, we changed the way the ZUI works. Originally, MindXpres used the CSS3 `transform: scale()` property to enlarge or reduce the whole view, giving the impression of zooming in or out. This is an obvious approach, simple in its execution and quite foolproof. However, the downside is that you can't zoom in very much, because currently browsers do not leverage the advantage of vector graphics and fonts even if you do use them, and obviously raster-based content doesn't scale much anyway. Instead, browsers render the content at its initial scale, and then treat the

result as one big image when scaled or otherwise transformed afterwards. This means you get extremely pixelated content when zooming in too much.

Through some refactoring, we were able to change this to use the `transform: translateZ()` property instead, along with the `transform: perspective()`⁵ and the `transform-style: preserve-3d` properties. This means we're now effectively rendering the presentation in 3D, and moving our viewpoint around in the 3D space to center each slide or component in turn.

We believe this opens the door for even more visually impressive presentations, where content can be placed on different points along the Z-axis. This allows for example to place multiple slides behind one another, making for impressive zoom transitions between slides. The downside of this is that the overview may not always show all content, as some content can overlap, but we trust the author uses this feature wisely when manually adjusting the position of their slides. It may for example be useful to group slides together in this way, when there is too much content to show on one slide but creating a second, separate slide may break the flow of the presentation. In any case, our automated layout plugin won't currently generate slides positioned this way.

Plain HTML5

After investigating the inner workings of MindXpres and studying some example presentations, we were ready to start generating our own presentations based on our component tree. This meant every possible component would have to be written out as valid HTML, with the necessary attributes for each generated tag and with any child components enclosed. Since our component tree nodes are nested the way the final HTML should be nested, this was not a problem.

Our implementation currently includes `compile()` methods on every component object, which is consistent and easy to understand, but which might not be the best way to implement this depending on future goals. We have to walk through the entire tree in any case, so performance will always be $O(n)$ at best. The current implementation has the advantage of extensibility, where new components can easily be added and it is immediately obvious to any new developer how these new components should generate their equivalent HTML code. However, for replacing the output with a different format

⁵ Not to be confused with the `perspective:`

number

property, which yields different results

it would be better if this functionality was separated from the components and gathered in a distinct `Writer` class instead. Switching formats would then be as easy as dropping in a new `Writer` class that generates a different format. Since we initially did not expect to switch outputs, and because we started implementing one component and then added components as we needed them, the current implementation — focussed on simplifying addition of new components — was easier for us to work with. Perhaps the refactoring of this implementation is an option for future work.

MindXpres XML

Generating MindXpres XML should be simpler than the HTML5 output, although in the end there is probably not much difference. We won't have to generate unique ID's for every component, and we won't have to generate the preamble content (which includes the MindXpres library itself), but the structure and content should remain pretty much the same. Instead of `<div>` tags we can now use MindXpres-specific tags. All that adds up to a much more easily readable result.

5.3 Creating layouts

5.3.1 Using constraints

5.3.2 Other ways

6. CONCLUSIONS AND FUTURE WORK

6.1 *Contribution*

Converting PowerPoint presentations into MindXpres is now possible and easy.

Automated layout has not been available in presentations until now. MindXpres is the first presentation system that doesn't require predefined templates, nor manual layout tweaking by the enduser, instead letting people focus on the content while the software takes care of the rest.

6.2 *Future Work*

Other formats: extend the convertor tool to convert Keynote, prezzi, pptx, ...

Other formats: refactor the convertor tool to separate output from the component classes, concentrating it in a Writer class that implements a Writer interface, so that the Writer class can easily be replaced with another one to generate a different format.

Integrate the convertor into an MindXpres editor

Improve the automated layout algorithm beyond constraints using learning AI, training it on good/bad layouts, neural network...

Implement extreme zooming into MindXpres using the CSS3 perspective property along with translation along the z-axis to utilize the full power of the 3d space.

BIBLIOGRAPHY

- I. Parker, “Absolute PowerPoint: Can a software package edit our thoughts?” *The New Yorker*, pp. 76–87, May 2001. [Online]. Available: http://www.newyorker.com/archive/2001/05/28/010528fa_fact_parker?currentPage=all
- S. M. Drucker, G. Petschnigg, and M. Agrawala, “Comparing and managing multiple versions of slide presentations,” in *Proceedings of the 19th Annual ACM Symposium on User Interface Software and Technology*, ser. UIST ’06. New York, NY, USA: ACM, 2006, pp. 47–56. [Online]. Available: <http://doi.acm.org/10.1145/1166253.1166263>
- G. Bajaj. (2013, Feb.) 30 million PowerPoint presentations? [Online]. Available: <http://blog.indezine.com/2013/02/30-million-powerpoint-presentations.html>
- E. R. Tufte. (2005, Sep.) PowerPoint does rocket science — and better techniques for technical reports. [Online]. Available: http://cgvr.cs.uni-bremen.de/miscellaneous/presentation_tips/edward_tufte/ppt_does_rocket_science.shtml
- R. Roels and B. Signer, “Mindxpres: An extensible content-driven cross-media presentation platform,” in *Web Information Systems Engineering — WISE 2014*, ser. Information Systems and Applications, incl. Internet/Web, and HCI, vol. 8787. Springer International Publishing, 2014, pp. 215–230. [Online]. Available: <http://link.springer.com/book/10.1007%2F978-3-319-11746-1>
- E. R. Tufte, *The Cognitive Style of PowerPoint: Pitching Out Corrupts Within*. Graphics Press, 2003.
- D. K. Farkas, “Toward a better understanding of PowerPoint deck design,” *Information Design Journal*, vol. 14, no. 2, pp. 162–171, 2006.
- T. H. Nelson, “The heart of connection: Hypermedia unified by transclusion,” *Commun. ACM*, vol. 38, no. 8, pp. 31–33, Aug. 1995. [Online]. Available: <http://doi.acm.org/10.1145/208344.208353>

- A. Holzinger, M. Kickmeier-Rust, and D. Albert, "Dynamic media in computer science education; content complexity and learning performance: Is less more?" *Journal of Educational Technology & Society*, vol. 11, no. 1, pp. 279–290, Jan. 2008.
- S. Lok and S. Feiner, "A survey of automated layout techniques for information presentations," 2001.
- C. Adams, "PowerPoint, habits of mind, and classroom culture," *Journal of Curriculum Studies*, vol. 38, no. 4, pp. 389–411, 2006.
- A. Gross and J. Harmon, "The structure of PowerPoint presentations: The art of grasping things whole," *IEEE Transactions on Professional Communication*, vol. 52, no. 2, pp. 121–137, Jun. 2009.
- L. Good and B. B. Bederson, "Zoomable user interfaces as a medium for slide show presentations," *Information Visualization*, vol. 1, no. 1, pp. 35–49, Mar. 2002. [Online]. Available: <http://dx.doi.org/10.1057/palgrave/ivs/9500004>
- L. Lichtschlag, T. Karrer, and J. Borchers, "Fly: A tool to author planar presentations," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, ser. CHI '09. New York, NY, USA: ACM, 2009, pp. 547–556. [Online]. Available: <http://doi.acm.org/10.1145/1518701.1518786>
- J. Lanir, K. S. Booth, and A. Tang, "Multipresenter: A presentation system for (very) large display surfaces," in *Proceedings of the 16th ACM International Conference on Multimedia*, ser. MM '08. New York, NY, USA: ACM, 2008, pp. 519–528. [Online]. Available: <http://doi.acm.org/10.1145/1459359.1459428>
- J.-C. Chen, W.-T. Chu, J.-H. Kuo, C.-Y. Weng, and J.-L. Wu, "Tiling slideshow," in *Proceedings of the 14th Annual ACM International Conference on Multimedia*, ser. MULTIMEDIA '06. New York, NY, USA: ACM, 2006, pp. 25–34. [Online]. Available: <http://doi.acm.org/10.1145/1180639.1180653>
- B. Signer and M. C. Norrie, "Paperpoint: A paper-based presentation and interactive paper prototyping tool," in *Proceedings of the 1st International Conference on Tangible and Embedded Interaction*, ser. TEI '07. New York, NY, USA: ACM, 2007, pp. 57–64. [Online]. Available: <http://doi.acm.org/10.1145/1226969.1226981>

- L. Nelson, S. Ichimura, E. R. Pedersen, and L. Adams, "Palette: A paper interface for giving presentations," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, ser. CHI '99. New York, NY, USA: ACM, 1999, pp. 354–361. [Online]. Available: <http://doi.acm.org/10.1145/302979.303109>
- R. Spicer, Y.-R. Lin, A. Kelliher, and H. Sundaram, "Nextslideplease: Authoring and delivering agile multimedia presentations," *ACM Trans. Multimedia Comput. Commun. Appl.*, vol. 8, no. 4, pp. 53:1–53:20, Nov. 2012. [Online]. Available: <http://doi.acm.org/10.1145/2379790.2379795>
- D. Edge, J. Savage, and K. Yatani, "Hyperslides: Dynamic presentation prototyping," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, ser. CHI '13. New York, NY, USA: ACM, 2013, pp. 671–680. [Online]. Available: <http://doi.acm.org/10.1145/2470654.2470749>
- P. Garcia, "Retooling PowerPoint for hypermedia authoring," in *Proceedings of Society for Information Technology & Teacher Education International Conference 2004*, R. Ferdig, C. Crawford, R. Carlsen, N. Davis, J. Price, R. Weber, and D. A. Willis, Eds. Atlanta, GA, USA: Association for the Advancement of Computing in Education (AACE), 2004, pp. 4098–4099. [Online]. Available: <http://www.editlib.org/p/14625>
- D. Austin, "Beginnings of PowerPoint: A personal technical story," 2009. [Online]. Available: <http://www.computerhistory.org/collections/catalog/102745695>
- B. Parks, "Death to PowerPoint," 2012.
- C. Sü and B. Freitag, "LMML — the learning material markup language framework," in *Proceedings of Workshop ICL*, Sep. 2002.
- J. Fisler and S. Bleisch, "eLML, the elesson markup language: Developing sustainable e-learning content using an open source XML framework," in *WEBIST 2006 — International Conference on Web Information Systems and Technologies*, Apr. 2006.
- J. H. Canós-Cerdá, M. I. Marante, and M. Llavador, "SliDL: A slide digital library supporting content reuse in presentations," in *ECDL 2010*, Sep. 2010.

- K. Verbert, X. Ochoa, and E. Duval, "The ALOCOM framework: Towards scalable content reuse," *Journal of Digital Information*, vol. 9, no. 1, pp. 1–24, Jan. 2008.
- D. Raggett, "Slidy — a web based alternative to Microsoft PowerPoint," in *XTech 2006*, 2006.
- M. D. Bush and J. D. Mott, "The transformation of learning with technology: Learner-centricity, content and tool malleability, and network effects," *Educational Technology*, vol. 49, no. 2, pp. 3–20, Mar. 2009.
- H. Haller and A. Abecker, "imapping: A zooming user interface approach for personal and semantic knowledge management," *SIGWEB Newsl.*, no. Autumn, pp. 4:1–4:10, Sep. 2010. [Online]. Available: <http://doi.acm.org/10.1145/1850770.1836295>
- R. J. Dufresne, W. J. Gerace, W. J. Leonard, J. P. Mestre, and L. Wenk, "Classtalk: A classroom communication system for active learning," *Journal of Computing in Higher Education*, vol. 7, no. 2, pp. 3–47, 1996.
- B. Signer and M. C. Norrie, "Interactive paper: Past, present and future," in *PaperComp 2010*, Sep. 2010.
- E. Reuss, B. Signer, and M. C. Norrie, "PowerPoint multimedia presentations in computer science education: What do users need?" in *USAB 2008*, 2008.
- R. Roels, C. Vermeulen, and B. Signer, "A unified communication platform for enriching and enhancing presentations with active learning components," in *ICALT 2014*, Jul. 2014, pp. 131–135.
- B. Signer and M. C. Norrie, "As we may link: A general metamodel for hypermedia systems," in *Proceedings of the 26th International Conference on Conceptual Modeling*, ser. ER'07. Berlin, Heidelberg: Springer-Verlag, 2007, pp. 359–374. [Online]. Available: <http://dl.acm.org/citation.cfm?id=1784489.1784522>
- N. Hurst, W. Li, and K. Marriott, "Review of automatic document formatting," in *Proceedings of the 9th ACM Symposium on Document Engineering*, ser. DocEng '09. New York, NY, USA: ACM, 2009, pp. 99–108. [Online]. Available: <http://doi.acm.org/10.1145/1600193.1600217>
- G. Di Battista, P. Eades, R. Tamassia, and I. G. Tollis, *Graph Drawing: Algorithms for the Visualization of Graphs*, 1st ed. Upper Saddle River, NJ, USA: Prentice Hall PTR, 1998.

- G. Di Battista, W.-P. Liu, and I. Rival, "Bipartite graphs, upward drawings, and planarity," *Inf. Process. Lett.*, vol. 36, no. 6, pp. 317–322, Dec. 1990. [Online]. Available: [http://dx.doi.org/10.1016/0020-0190\(90\)90045-Y](http://dx.doi.org/10.1016/0020-0190(90)90045-Y)
- F. Shahrokhi, L. A. Szkely, O. Skora, and I. Vrt'o, "Drawings of graphs on surfaces with few crossings," *Algorithmica*, vol. 16, no. 1, pp. 118–131, 1996, *algorithmica*, 16(1):118131, July 1996. [Online]. Available: <http://dx.doi.org/10.1007/BF02086611>
- T. C. Hu and E. S. Kuh, "VLSI circuit layout: Theory and design," 1985, IEEE, USA, 1985.
- T. Lengauer, *Combinatorial Algorithms for Integrated Circuit Layout*. New York, NY, USA: John Wiley & Sons, Inc., 1990.
- V. J. Milenkovic, K. M. Daniels, and Z. Li, "Automatic marker making," 1991, in T. Shermer, editor, *Proceedings of the Third Canadian Conference on Computational Geometry*, pages 243246, August 610 1991.
- M. Hofri, "Two-dimensional packing: Expected performance of simple level algorithms," *Information and Control*, 45:117, 1980., no. 45, pp. 1–17, 1980. [Online]. Available: <http://www.sciencedirect.com/science/article/pii/S0019995880908177>
- J. Scahill and J. Begley. (2015, Mar.) The CIA campaign to steal Apple's secrets. [Online]. Available: <https://firstlook.org/theintercept/2015/03/10/ispy-cia-campaign-steal-apples-secrets>
- J. Gosling and H. McGilton, "The Java language environment," 1995. [Online]. Available: http://www.stroustrup.com/1995_Java_whitepaper.pdf
- E. Sundaram, "Excelling in excel with Java," 2004.